

OMSCS GEORGIA TECH

Spanning Tree

CS 6250

Summer 2022

Copyright 2021

Georgia Institute of Technology

All rights reserved.

This is solely to be used for current CS6250 students. Any public posting of the material contained within is strictly forbidden by the Honor code.

Spanning Tree

Table of Contents

PROJECT GOAL.....	2
Part 1: Setup	2
Part 2: Files Layout.....	2
Part 3: TODOs	3
Part 4: Testing and Debugging	6
Part 5: Assumptions and Clarifications	6
What to Turn In	7
What you can and cannot share.....	8
Rubric	9

PROJECT GOAL

In the lectures, you learned about Spanning Trees (see Canvas->Modules->Lesson 1-> Looping Problem in Bridges and the Spanning Tree Algorithm), which can be used to prevent forwarding loops on a layer 2 network (see the Wikipedia entry on [Spanning Trees](#)). In this project, you will develop a simplified distributed version of the [Spanning Tree Protocol](#) that can be run on an arbitrary layer 2 topology. With this project, we will be simulating the communications between switches until they converge on a single solution, and then output the final spanning tree to a file.

This document covers the following: 1. Project Files Layout; 2. Project Outline: TODOs; 3. Testing and Debugging; 4. Assumptions and Clarifications; 5. Submission; 6. Grading; 7. Honor Code/Academic Integrity.

Part 1: Setup

Download the project from Canvas to the course VM and unzip it. Alternatively, you can do this project on your host system if it has Python 3.8 or newer installed. The project does not have many dependencies outside of Python. You must be sure that your submission runs properly in Gradescope. Gradescope is the environment where your project will be graded.

Part 2: Files Layout

There are many files in the SpanningTree directory, but you should only modify *Switch.py*, which represents a layer 2 switch that implements our simple Spanning Tree Protocol. You will implement the functionality of the Spanning Tree Protocol to generate a Spanning Tree for each Switch.

The other files in the project skeleton are described below. DO NOT modify these files. All of your coding will be in Switch.py ONLY. However please study these files to understand the project better and to make decisions about your code in Switch.py.

- `Topology.py` - Represents a network topology of layer 2 switches. This class reads in the specified topology and arranges it into a data structure that your switch code can access.
- `StpSwitch.py` - A base class of the derived class you will code in Switch.py. The base class `StpSwitch.py` abstracts certain implementation details to simplify your tasks.
- `Message.py` - This class represents a simple message format you will use to communicate between switches, similar to what was described in the course lectures. Specifically, you will create and send messages in Switch.py by declaring a message as

```
msg = Message(claimedRoot, distanceToRoot, originID,  
              destinationID, pathThrough)
```

and assigning the correct data to each input. Message format may not be changed. See the comments in `Message.py` for more information on the data in these variables.

- `run_spanning_tree.py` - A simple "main" file that loads a topology file (see `XXXTopo.py` below), uses that data to create a Topology object containing Switches, and starts the simulation.
- `XXXTopo.py`, etc. - These are topology files that you will pass as input to the `run_spanning_tree.py` file.
- `sample_output.txt` - Example of a valid output file for `Sample.py` as described in the comments in `Switch.py`.

Part 3: TODOs

This is an outline of the code you must implement in `Switch.py` with *suggestions* for implementation. Your implementation must adhere to the "spirit of the project" and satisfy the labeled TODO sections in the project code per the pre-existing comments. We also provide a walkthrough video that past students have found helpful. Please keep an eye out for that link on Ed Discussion.

A. Decide on the data structure that you will use to keep track of the spanning tree.

1. The collection of active links across all switches is the resultant spanning tree.
2. The data structure may be variable(s) needed to track each switch's own view of the tree. **A switch only has access to its member variables. A switch may not access its neighbor's information directly – to learn information from a neighbor, the neighbor must send a message.**
3. This is a distributed algorithm. The switch can only communicate with its neighbors. It does not have an overall view of the spanning tree, or of the topology as a whole.
4. An example data structure would include, at a minimum:
 - a. a variable to store the switch ID that this switch currently sees as the *root*,
 - b. a variable to store the *distance* to the switch's root,
 - c. a [list](#) or other datatype that stores the "*active links*" (i.e., the links to neighbors that should be drawn in the spanning tree).
 - d. a variable to keep track of which neighbor it goes through to get to the root. (A switch should only go through one neighbor, if any, to get to the root.)
5. More variables may be helpful to track data needed to build the spanning tree and will depend on your specific implementation.

6. It is important to create a data structure in the correct place in Python (and in most object-oriented programming languages). If you create it inside a method, every time the method is called it will be recreated. You should create a class object in the class constructor so that the data stored in the object exists for the life of the class instance that is created by `Topology.py`. For example `self.mylist = []` in the constructor should create an empty list data structure and act as instance variable. But if `mylist` were instantiated in, say, `process_messages`, then it will be created every time the method is called. This could be useful in how you track which links are active to certain neighbors for any given switch.

B. Implement sending initial messages to neighbors of the switch.

1. Your implementation of `send_initial_messages()` in `Switch.py` will be called in `Topology.py` for each switch in the topology before any other messages are processed and/or sent.
2. See code comments in `Message.py`, `Topology.py`, and `StpSwitch.py` for details on message format, message creation, and how to send messages between switches.
 - a. `pathThrough` is a Boolean, not an int.
 - b. In a message, *pathThrough* is TRUE if the message-sending switch goes through the message-receiving switch in order to reach *claimedRoot*. *pathThrough* is FALSE if the message-sending switch does not go through the message-receiving switch in order to reach *claimedRoot*.
3. Initially, each switch thinks it is the root of the spanning tree.

C. Implement processing a message from an immediate neighbor.

1. For each message a switch receives, the switch will need to:
 - a. **Determine whether an update to the switch's root information is necessary and update accordingly.**
 - I. The switch should update the *root* stored in its data structure if it receives a message with a lower *claimedRoot*.
 - II. The switch should update the *distance* stored in its data structure if a) the switch updates the *root*, or b) there is a shorter path to the same root.

1 - 2, 1 - 3	1 - 3, 1 - 2
2 - 1, 2 - 4	2 - 4, 2 - 1
3 - 1	3 - 1
4 - 2	4 - 2

Part 4: Testing and Debugging

To run your code on a specific topology (SimpleLoopTopo.py in this case) and output the results to a text file (out.txt in this case), execute the following command:

```
python run_spanning_tree.py SimpleLoopTopo out.txt
```

NOTE: “SimpleLoopTopo” is not a typo in the example command – don’t include the .py extension.

We have included several topologies with correct solutions for you to test your code against. You can (and are encouraged to) create more topologies and test suites with output files and share them on Ed Discussion. There will be a post where students can share these files.


You will only be submitting `Switch.py` – your implementation must be confined to modifications of that file. We recommend testing your submission against a clean copy of the rest of the project files prior to submission.

We encourage adding print statements to facilitate debugging during the development process, but they should be removed/commented out prior to submission.

Part 5: Assumptions and Clarifications

You may assume the following:

- A. **All switch IDs are positive integers, and distinct.**
 1. These integers do not have to be consecutive.
 2. They will not always start at 1.
 3. There is no maximum value beyond language (Python) limitations (which your code does not need to check for).
- B. **Tie breakers:** If there are two paths of equal distance to the same root, the switch should choose the path through the neighbor with the lowest switch ID.
 1. Example: switch 5 has two paths to root switch 1, through switch 3 and switch 2. Each path is 2 hops in length. Switch 5 should select switch 2 as the path to the root and disable forwarding on the link to switch 3.



- C. **There is a single distinct solution spanning tree for each topology.** This is guaranteed by the first two assumptions.
- D. **All switches in the network will be connected to at least one other switch, and all switches are able to reach every other switch.** It will always be possible to form a tree that spans the entire topology.
- E. **There will be only 1 link between each pair of directly connected switches.** You do not need to consider how STP should behave with redundant links.
- F. **The topology given at the start will be the final topology.** The topology will not change while your code is running (i.e., adding a switch, severing a connection, etc.)
- G. **A switch may always communicate with its neighbors.** When a switch treats a link as inactive, the link can still be used during the simulation. “Inactive” simply means that the port/link will not be used for forwarding normal network traffic.
- H. **The solution implemented in `Switch.py` should terminate without intervention.** When there are no more messages in the queue to process, the simulation will log output and self-terminate.
- I. Your solution should not require any outside Python modules. Please do not import any other modules.

What to Turn In

To complete this project, submit **ONLY** your `Switch.py` file to Gradescope as a single file. Do **not** modify the name of `Switch.py` or else and you may receive a zero. You can make an unlimited number of submissions to Gradescope. Your last submission will be your grade unless you activate a different submission.

Before submission:

- a. **Make sure your logging format is correct.** Invalid format will result in autograder penalties.

- b. **Remove any print statements from your code before turning it in.** Print statements left in the simulation, especially for inefficient implementations, have drastic effects on runtime. Your submission must take less than 10 seconds to process a topology used in grading. If print statements in your code adversely affect the grading process, your work will not receive full credit.
- c. **Make sure your Switch.py works in Gradescope!** Gradescope will give you immediate feedback with your grade, so we will not accept re-grade requests related to incorrect submissions.
- d. **Make sure your Switch.py has the proper Linux-style line endings!** Specifically, if you are editing the files in Windows, make sure the line endings remain Linux-style, with just a LF at the end of lines. Windows may try to put CRLF at the end of lines.

After submission:

- e. **Make sure your submission uploaded correctly.** Late submissions will not be accepted.
- f. **Your grade in Gradescope will be your grade for this project**, with some caveats:
 - a. Any Honor Code violations will result in a 0 and be referred to OSI.
 - b. Any hard-coded solutions that violate the “spirit of the project” will be regraded manually
 - c. Any attempt to bypass or distort the autograder will result in a 0 and may be referred to OSI.

What you can and cannot share

Honor Code/Academic Integrity: Do **not** share any code from `Switch.py` with your fellow students, on Ed Discussion, or publicly in any form. You **may** share log files for any topology, and you may also share any code you write that will *not be turned in*, such as new topologies or other testing code.

In past semesters, the most trouble we have had with students not abiding by the honor code was in the Spanning Tree Project. All work must be your own, and consulting Spanning Tree Protocol solutions, even in another programming language or just “for reference”, are considered violations of the honor code. **Do not** reference solutions on Github! For more information see the Syllabus Definition of Plagiarism.

Start early, ask questions in Ed Discussion, and attend TA chats if needed. While this project is challenging, most of our students have succeeded with time and hard work and have a great sense of personal achievement with this project.

Rubric

10 pts	Correct Submission	For turning in the correct file with the correct name. You receive 10 free points for reading the instructions.
45 pts	Provided Topologies	For correct Spanning Tree results (log files) on the provided topologies.
45 pts	Unannounced Topologies	For correct Spanning Tree results (log files) on three topologies that you will not see in advance. They are slightly more complex than the provided ones and may test for corner cases.

Grading Note: Partial credit is **not** available for individual topology outputs/logs. The output must be fully correct to receive credit for that input topology – a single link’s discrepancy will result in a zero for that topology. Additionally, we will be using many topologies to test your project, including but not limited to the topologies we provide, and checking for corner cases not exhibited in the sample topologies provided.

The goal of this project is to implement a simplified version of a network protocol using a **distributed** algorithm. This means that your algorithm should be implemented at the network switch level. Each switch only knows its internal state, and the information passed to it via messages from its direct neighbors - the algorithm **must** be based on these messages.

The skeleton code we provide you runs a simulation of the larger network topology, and for the sake of simplicity, the `StpSwitch` class defines a link to the overall topology. This means it is possible using the provided code for one Switch to access another’s internal state. This goes against the spirit of the project and is not permitted. Additional detail is available in the comments of the skeleton code. **The autograder checks if submissions attempt to directly access `topolink` or `self.topology`. Submissions that attempt this will receive no credit** (If you have questions about whether your code is accessing data it should not, please ask on Ed Discussion or during office hours).